

Three Pillars with Zero Answers

A New Observability Scorecard

December 6, 2018

First, a Critique

The Conventional Wisdom

Observing microservices is hard

Google and Facebook solved this (right???)

They used **Metrics, Logging, and Distributed Tracing...**

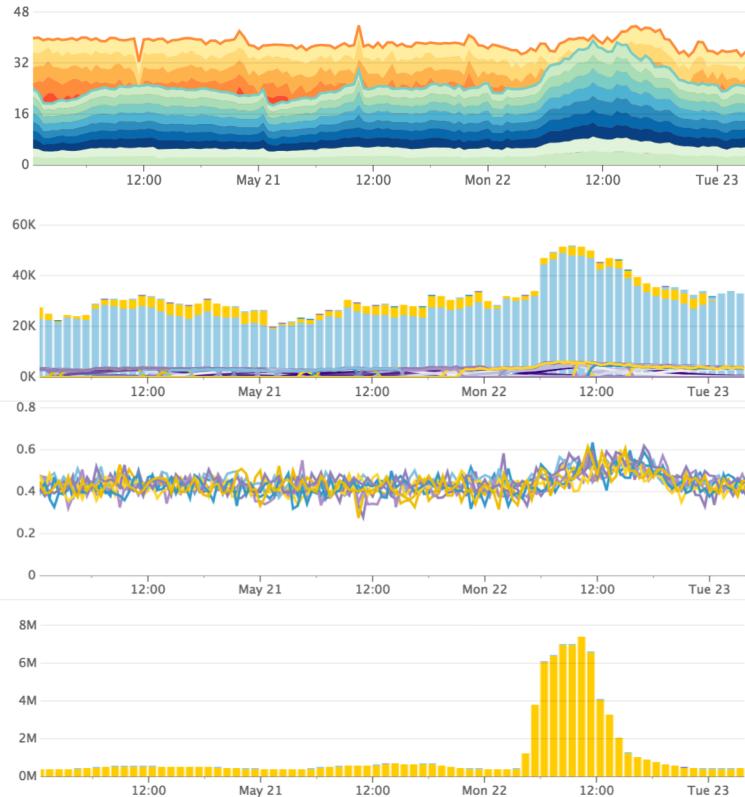
So we should, too.

The Three Pillars of Observability

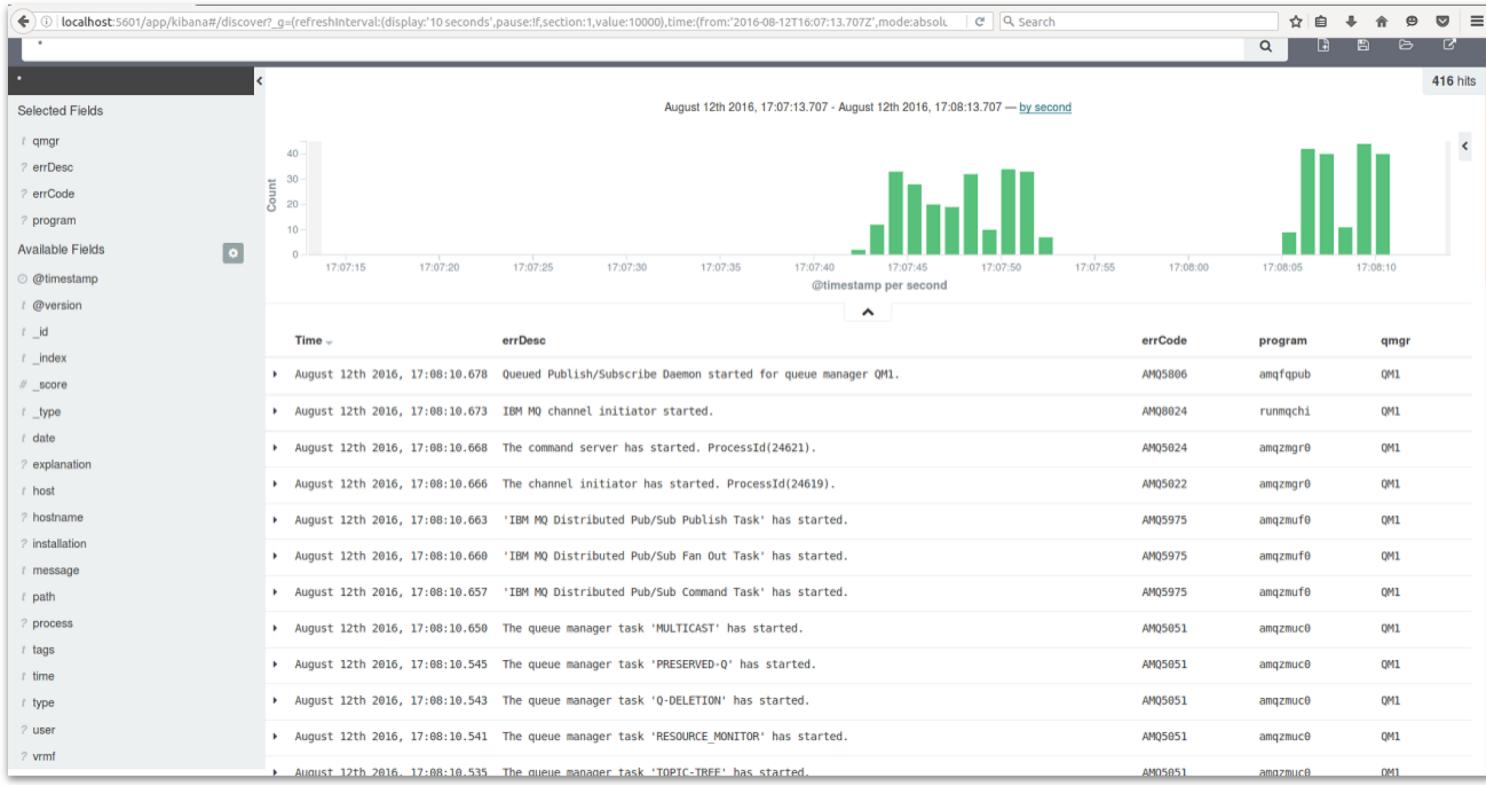
- Metrics
- Logging
- Distributed Tracing



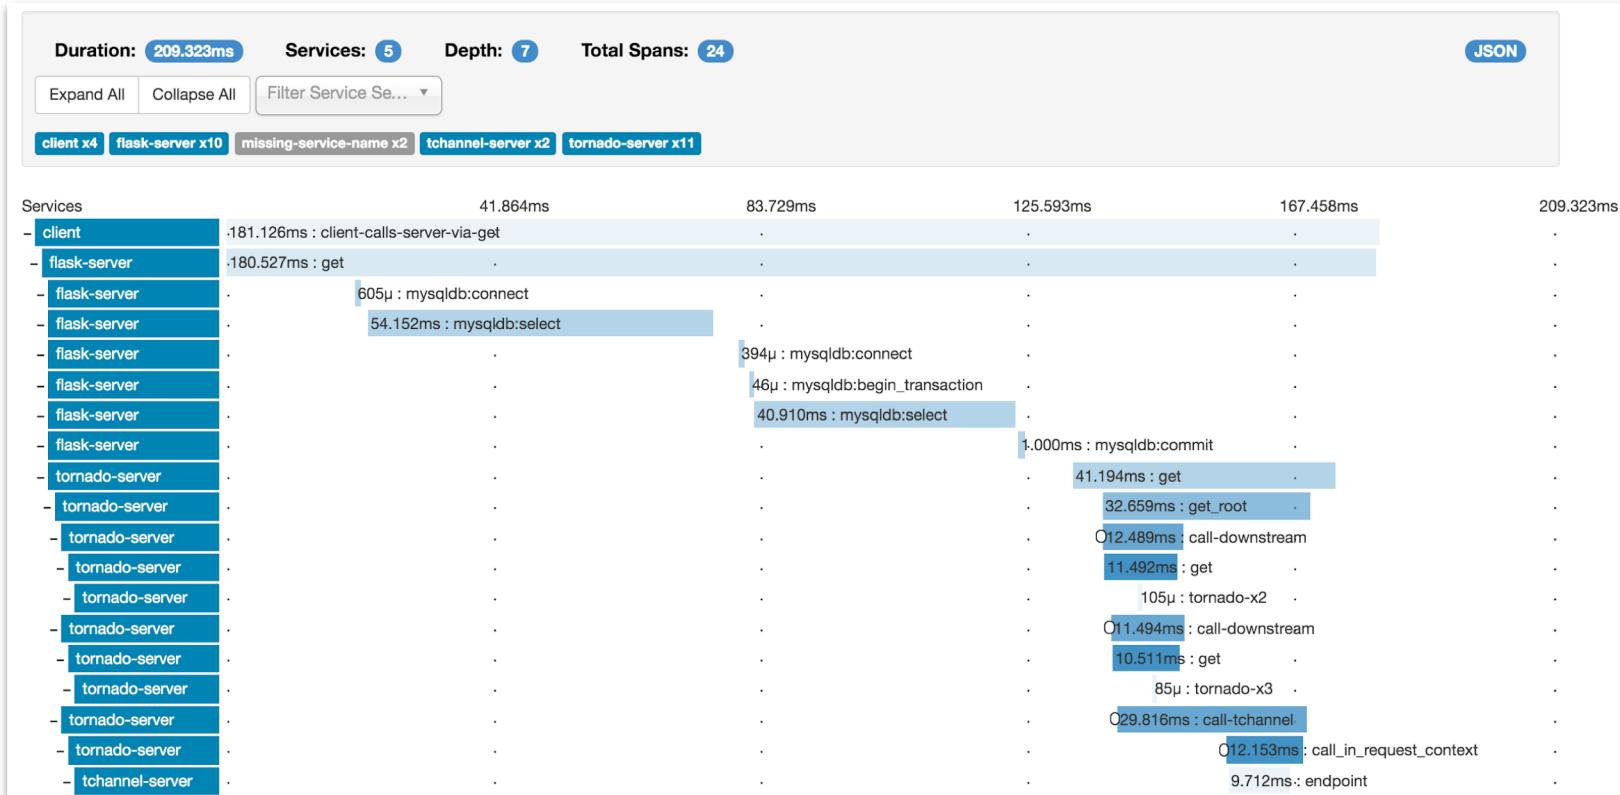
Metrics!



Logging!

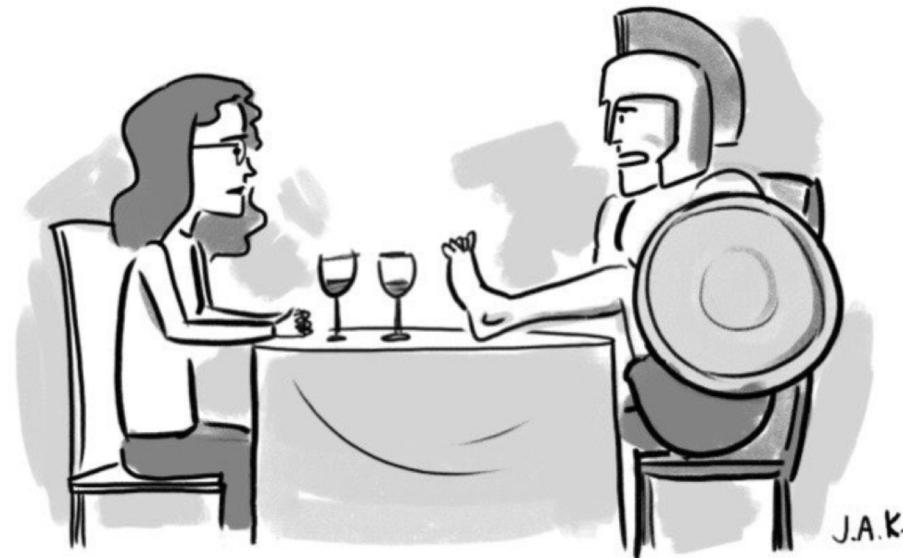


Tracing!





Fatal Flaws

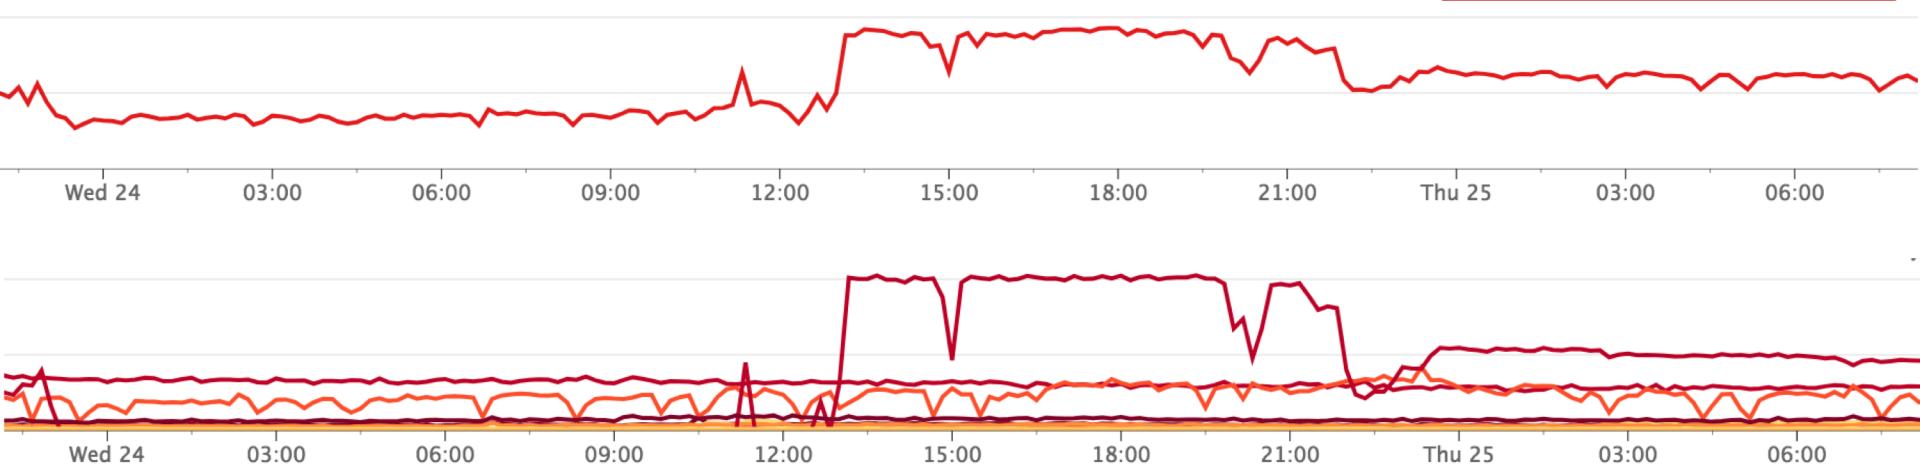


"I'm ready to be vulnerable."

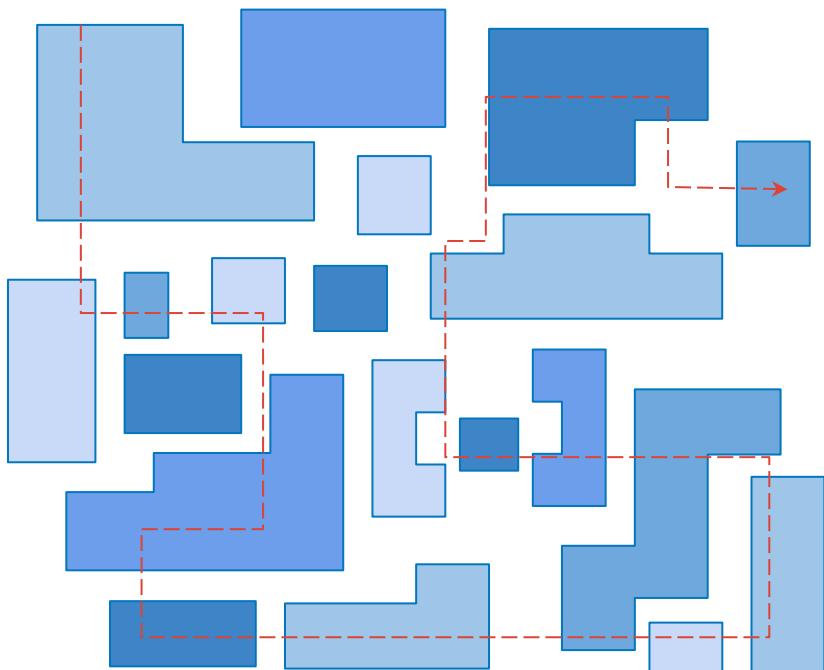
A word nobody knew in 2015...

Dimensions (aka “tags”) can explain variance
in timeseries data (aka “metrics”) ...

... but **cardinality**



Logging Data Volume: a reality check



- x transaction rate
- x all microservices
- x cost of net+storage
- x weeks of retention

way too much \$\$\$

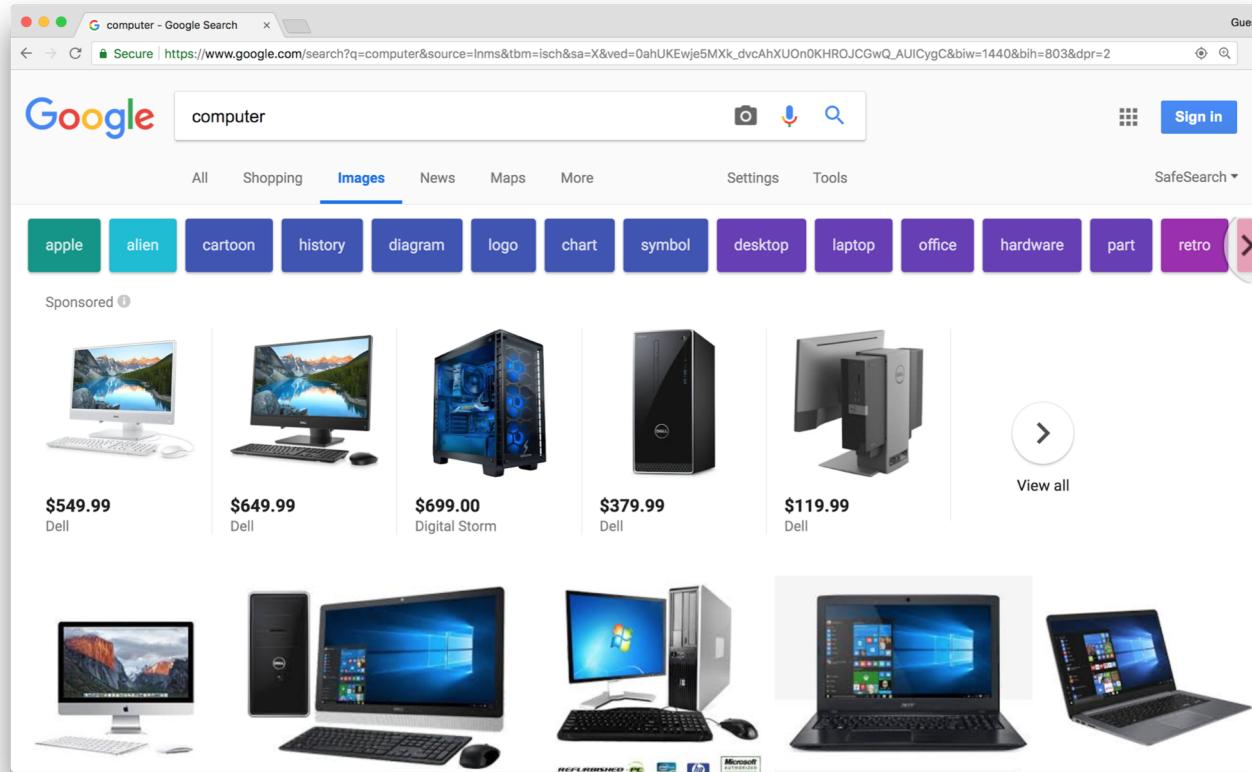
The Life of Transaction Data: Dapper

Stage	Overhead affects...	Retained
Instrumentation Executed	App	100.00%
Buffered within app process	App	000.10%
Flushed out of process	App	000.10%
Centralized regionally	Regional network + storage	000.10%
Centralized globally	WAN + storage	000.01%

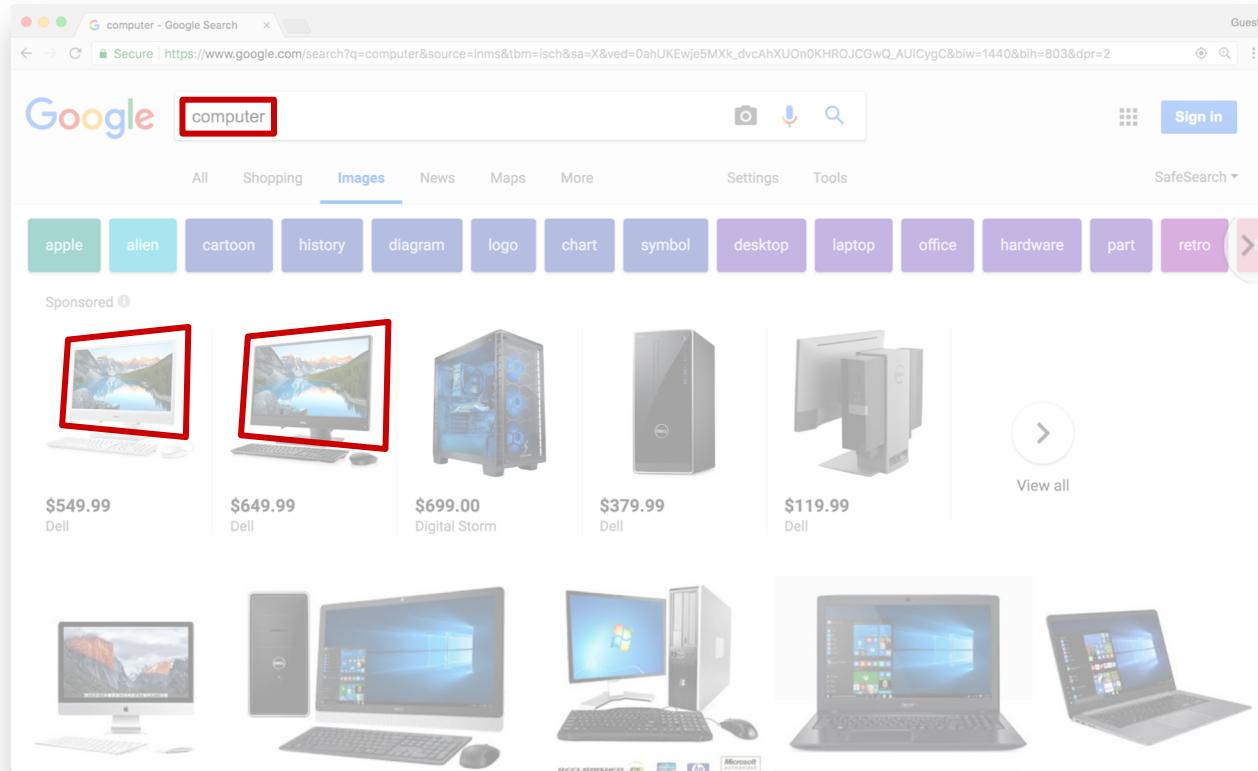
Fatal Flaws: A Review

	Logs	Metrics	Dist. Traces
TCO scales gracefully	—	✓	✓
Accounts for all data (i.e., unsampled)	✓	✓	—
Immune to cardinality	✓	—	✓

Data vs UI

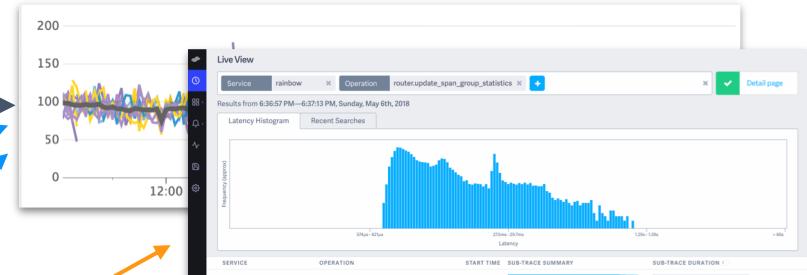


Data vs UI

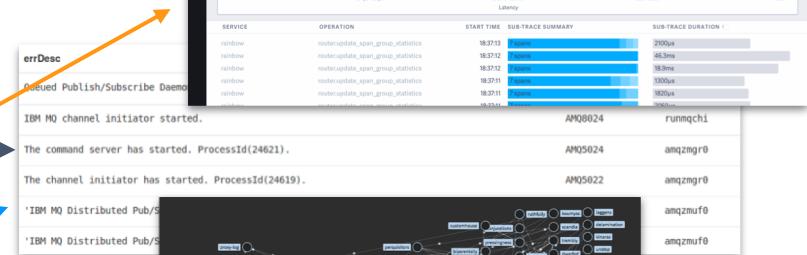


Data vs UI

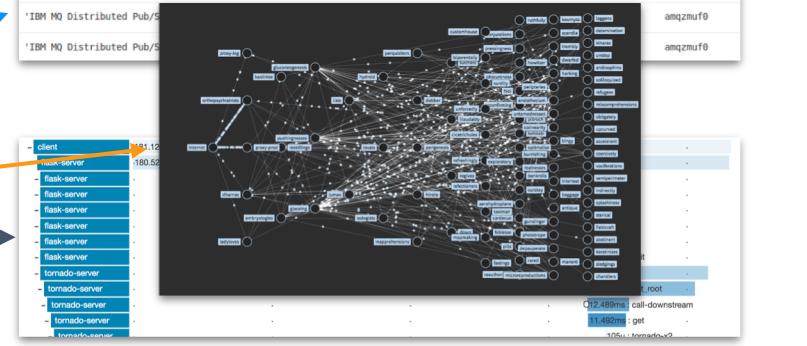
Metrics



Logs



Traces



Metrics, Logs, and Traces are
Just Data,

... not a feature or use case.

A New Scorecard for Observability

Observability: Quick Vocab Refresher

“SLI” = “Service Level Indicator”

TL;DR: An SLI is **an indicator of health** that a service’s **consumers** would care about.

... not an indicator of its inner workings

Observability: Two Fundamental **Goals**

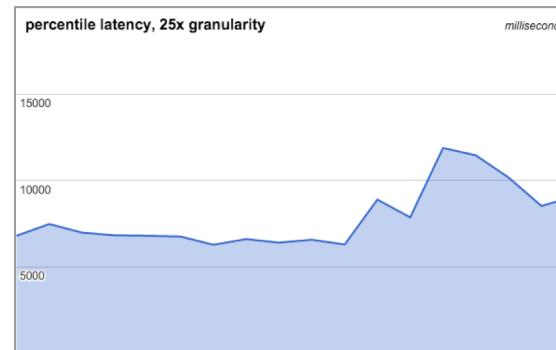
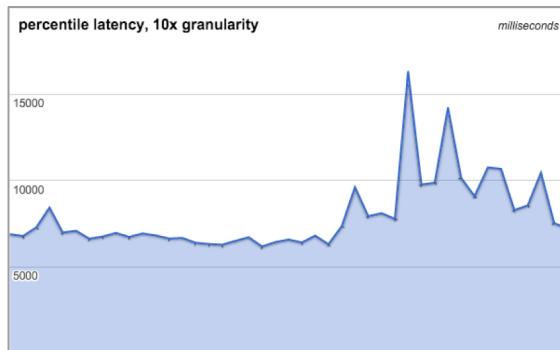
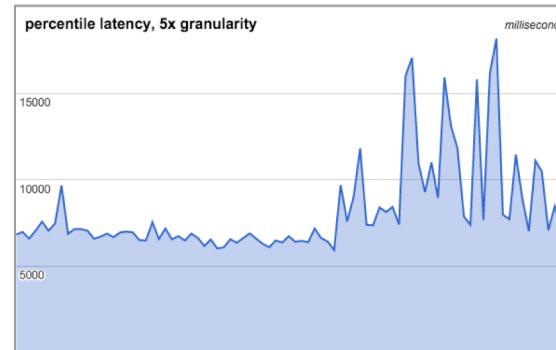
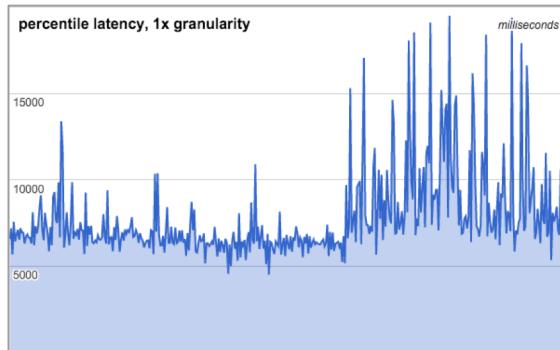
- Gradually improving an SLI
 - Rapidly restoring an SLI
-
- days, weeks, months...
- NOW!!!!

Reminder: “SLI” = “Service Level Indicator”

Observability: Two Fundamental **Activities**

1. Detection: perfect SLI capture
2. Refinement: reduce the search space

An interlude about stats frequency



Scorecard >> **Detection**

Specificity:

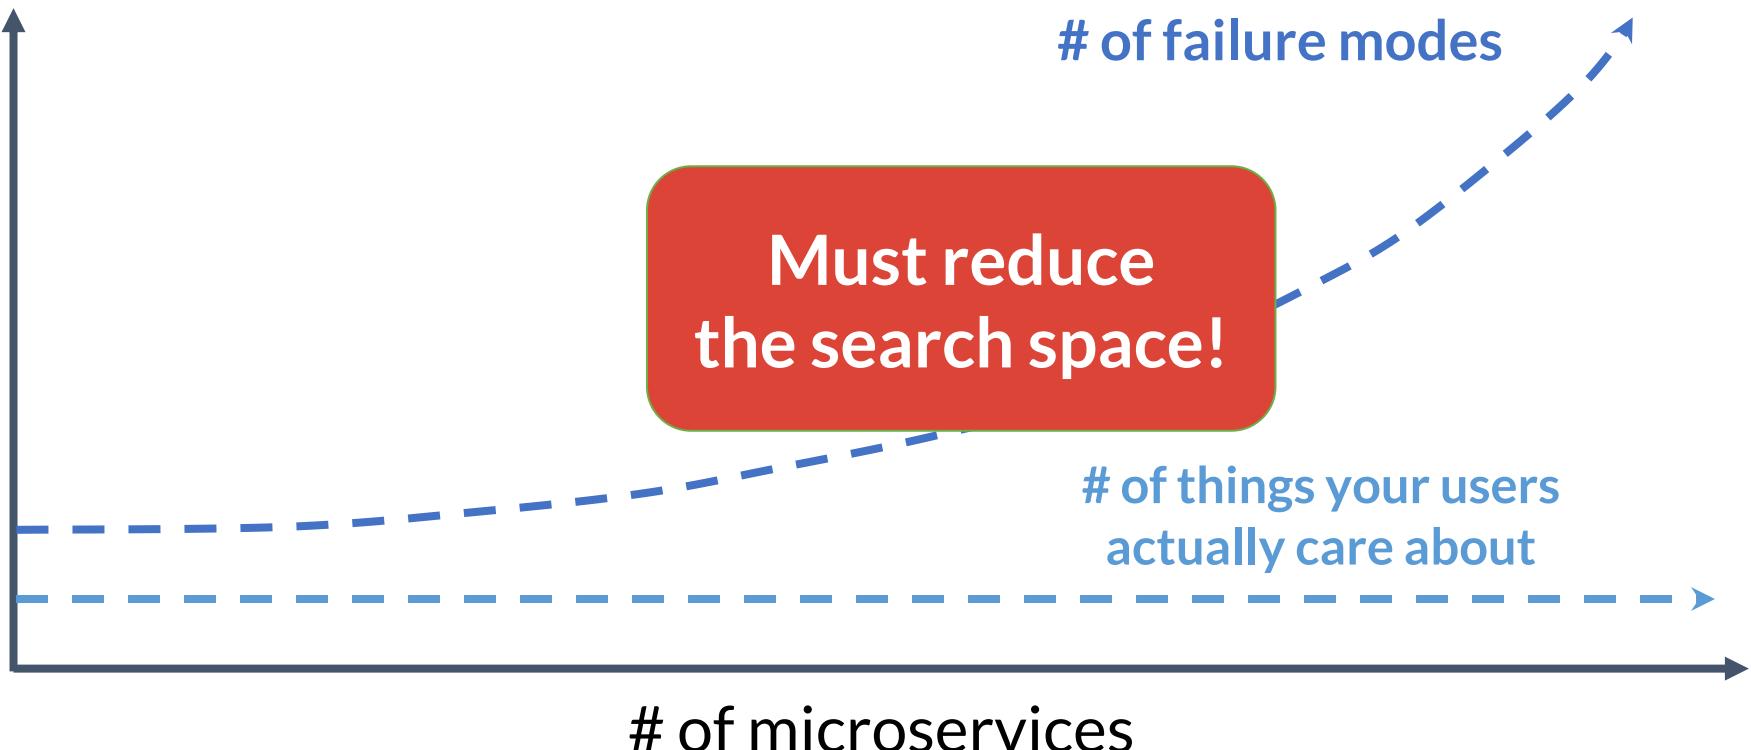
- Arbitrary dimensionality and cardinality
- Any layer of the stack, including mobile+web!

Fidelity:

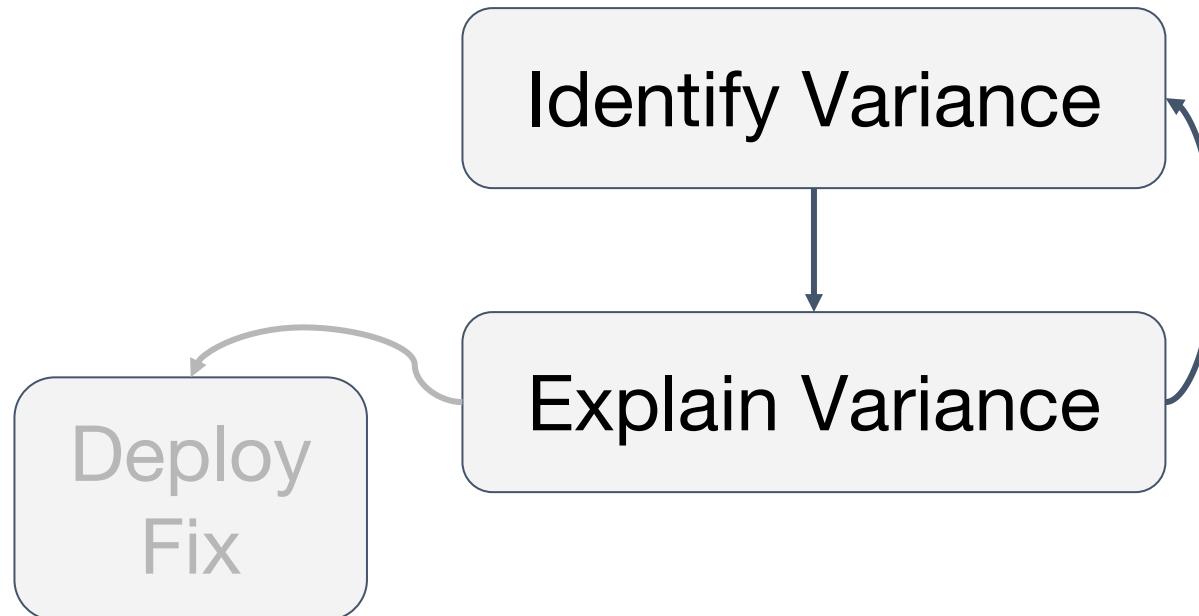
- Correct stats!!!
- High stats frequency (i.e., “beware smoothing”!)

Freshness: \leq 5 second lag

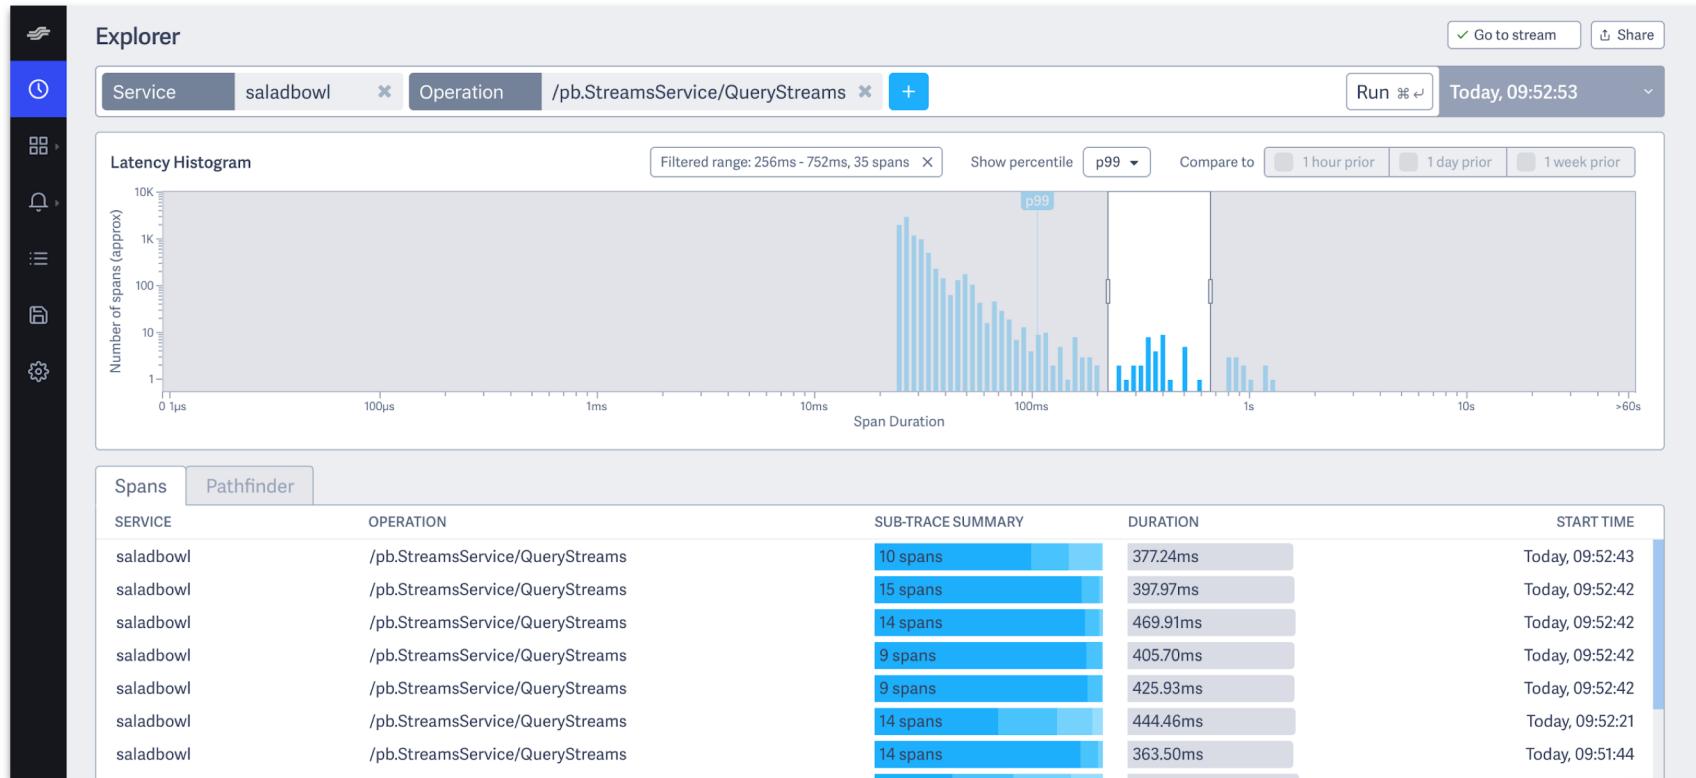
Scorecard >> Refinement



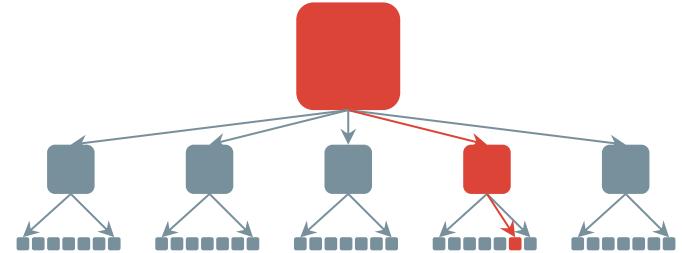
Scorecard >> **Refinement**



An interlude about variance and “p99”



Scorecard >> **Refinement**



Identifying Variance:

- Cardinality: understand which tag changed
- Robust stats: *histograms* (see prev slide)
- Data retention: always “Know What’s Normal”

Explaining variance:

- Correct stats!!!
- “Suppress the messengers” of microservice failures

Wrapping up...

(first, a hint at my perspective)

The Life of Transaction Data: Dapper

Stage	Overhead affects...	Retained
Instrumentation Executed	App	100.00%
Buffered within app process	App	000.10%
Flushed out of process	App	000.10%
Centralized regionally	Regional network + storage	000.10%
Centralized globally	WAN + storage	000.01%

The Life of Transaction Data: ~~Dapper~~ LightStep

Stage	Overhead affects...	Retained
Instrumentation Executed	App	100.00%
Buffered within app process	App	100.00%
Flushed out of process	App	100.00%
Centralized regionally	Regional network + storage	100.00%
Centralized globally	WAN + storage	on-demand

An Observability Scorecard

Detection

- Specificity: unlimited cardinality, across the entire stack
- Fidelity: correct stats, high stats frequency
- Freshness: ≤ 5 seconds

Refinement

- Identifying variance: unlimited cardinality, hi-fi histograms, data retention
- “Suppress the messengers”

Thank you!

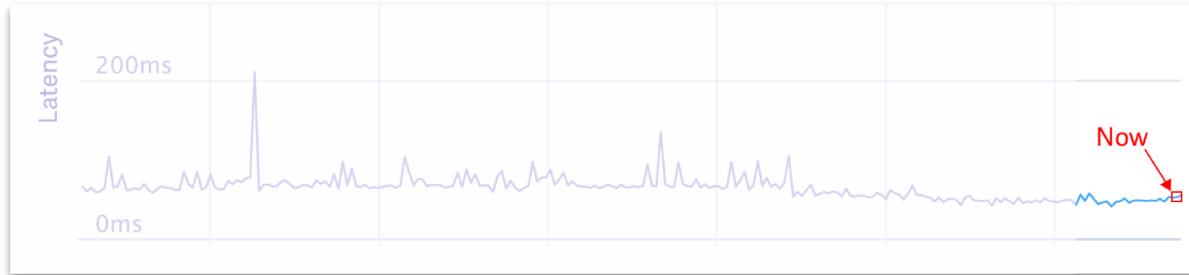
Ben Sigelman, Co-founder and CEO

twitter: @el_bhs

email: bhs@lightstep.com

Extra slides

Ideal Measurement: Robust



Ideal Measurement: High-Dimensional

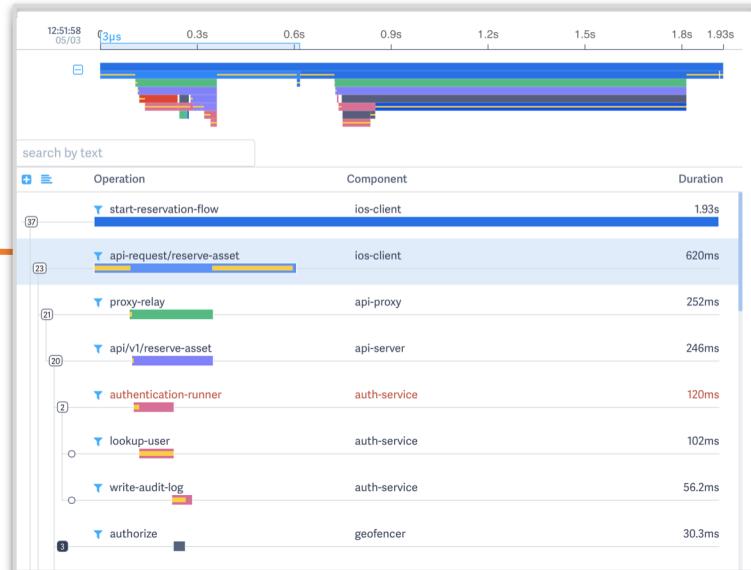
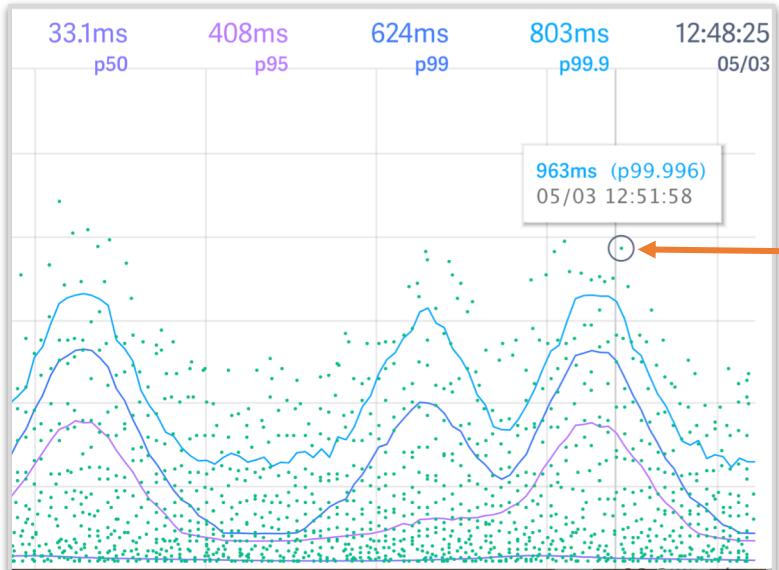


Ideal Refinement: Real-time

Must be able to test and eliminate hypotheses quickly

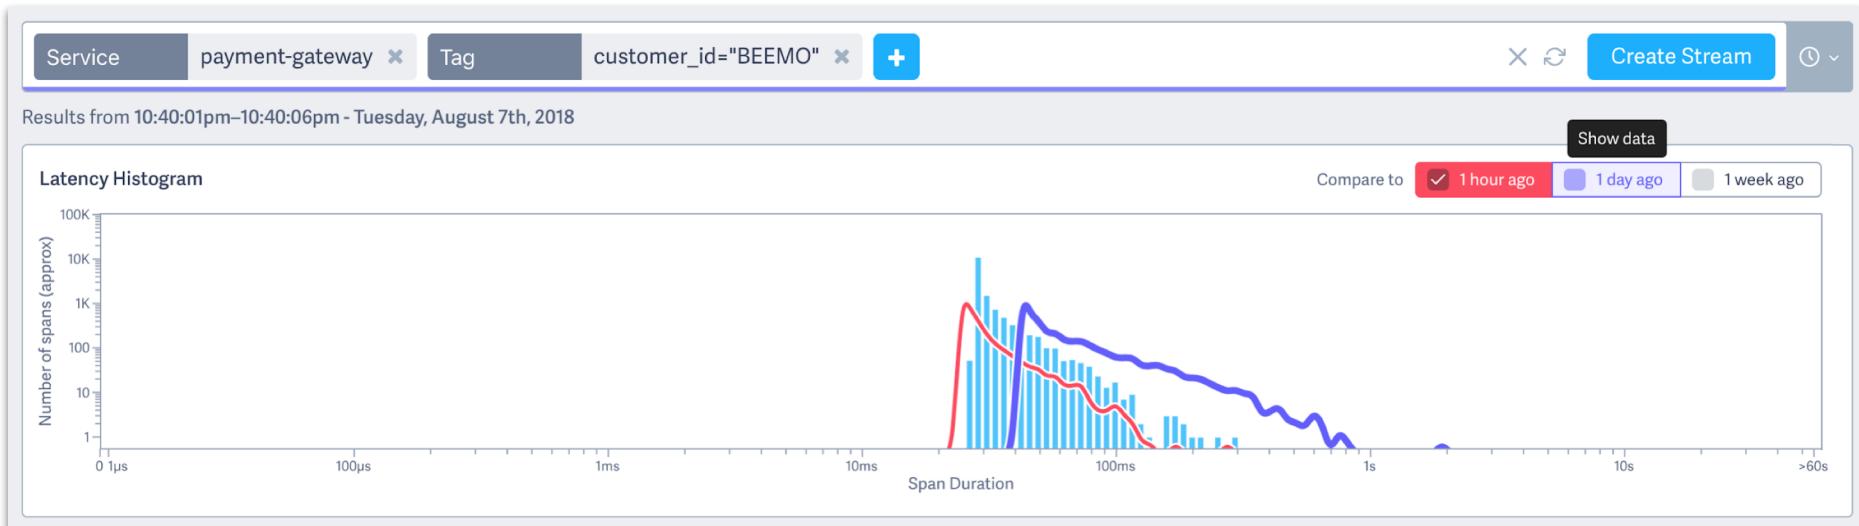
- Actual data must be $\leq 10\text{s}$ fresh
- UI / API latency must be very low

Ideal Refinement: Global



Ideal Refinement: Context-Rich

We can't expect humans to **know what's normal**



Thank you / Q&A